

HOMework 3

16824 VISUAL LEARNING AND RECOGNITION (SPRING 2025)

<https://piazza.com/cmu/spring2025/16824a/>

RELEASED: Thu, 20th Mar 2025

DUE: 11:59 PM ET, Thu, 10th Apr 2025

Instructor: Shubham Tulsiani

TAs: Hanzhe Hu, Homanga Bharadhwaj, Jianren Wang, Nikhil Keetha

START HERE: Instructions

- **Collaboration policy:** All are encouraged to work together BUT you must do your own work (code and write up). If you work with someone, please include their name in your write-up and cite any code that has been discussed. If we find highly identical write-ups or code or lack of proper accreditation of collaborators, we will take action according to strict university policies. See the [Academic Integrity Section](#) detailed in the initial lecture for more information.
- **Late Submission Policy:** There are a **total of 5** late days across all homework submissions. Submissions that use additional late days will incur a 10% penalty per late day.
- **Submitting your work:**
 - We will be using Gradescope (<https://gradescope.com/>) to submit the Problem Sets. Please use the provided template only. You do **not** need any additional packages and using them is **strongly discouraged**. Submissions must be written in LaTeX. All submissions not adhering to the template will not be graded and receive a zero.
 - **Deliverables:** Please submit all the `.py` files. Add all relevant plots and text answers in the boxes provided in this file. To include plots you can simply modify the already provided latex code. Submit the compiled `.pdf` report as well.

NOTE: Partial points will be given for implementing parts of the homework even if you don't get the mentioned accuracy as long as you include partial results in this pdf.

1 Image Captioning with Transformers (70 points)

We will be implementing the different pieces of a Transformer decoder ([Transformers](#)), and train it for image captioning on a subset of the [COCO dataset](#).


- **Setup:** Run the following command to extract COCO data, in the `transformer_captioning/datasets` folder: `./get_coco_captioning.sh`
- **Question:** Follow the instructions in the `README.md` file in the `transformer_captioning` folder to complete the implementation of the transformer decoder.
- **Deliverables:** After implementing all parts, use `run.py` for training the full model. The code will log plots to `plots`. Extract plots and paste them into the appropriate section below.
- **Expected results:** These are expected training losses after 100 epochs. Do not change the seed in `run.py`.
 - 2-heads, 2-layers, lr 1e-4: Final loss ≤ 1
 - 4-heads, 6-layers, lr 1e-4: Final loss ≤ 0.3
 - 4-heads, 6-layers, lr 1e-3: Final loss ≤ 0.05

1. Paste training loss plots for each of the three hyper-param configs


2-heads-2-layers-lr-1e-4: **0.776759**

4-heads-6-layers-lr-1e-4: **0.251522**


4-heads-6-layers-lr-1e-3: **0.043471**



(a) 2-heads-2-layers-lr-1e-4



(b) 4-heads-6-layers-lr-1e-4



(c) 4-heads-6-layers-lr-1e-3

2. Paste any three generated captioning samples from the training set with the three different settings. The provided code creates these plots at the end of training.

train
voman in a woman in the <UNK> the <UNK> tennis player with his hand <E
GT:<START> the woman in the <UNK> hold the box of donuts <END>



train
the woman in the <UNK> hold the box of donuts <END>
GT:<START> the woman in the <UNK> hold the box of donuts <END>



train
the woman in the <UNK> hold the box of donuts <END>
GT:<START> the woman in the <UNK> hold the box of donuts <END>



(a) Sample1: 2-heads-2-layers-lr-1e-4

(b) Sample2: 4-heads-6-layers-lr-1e-4

(c) Sample3: 4-heads-6-layers-lr-1e-3

3. Based on the observations of the three different settings, What would you change in the training procedure to get better validation performance? Why tweaking these hyper-parameters will lead to better performances?


Solution:

Improving validation performance can be achieved by optimizing crucial hyperparameters: learning rate, attention heads, and layers. Evidence suggests that raising model complexity (more attention heads and layers) helps capture intricate patterns in images. However, to maintain stability, adjust the learning rate appropriately, either by using a smaller rate or implementing schedulers like Cosine Annealing. To prevent overfitting in complex models, apply regularization techniques such as dropout and weight decay via optimizers like AdamW.


2 Classification with Vision Transformers (30 points)

We will use the transformer you implemented in the previous part to implement a Vision Transformer (ViT), for classification on CIFAR10.

- **Question:** Follow the instructions in the `README.md` file in the `vit_classification` folder. You are encouraged to reuse code from the previous question.
- **Deliverables:** Run training using `run.py` for training the full model. The code will log plots `acc_out.png` (train and test accuracy) and `loss_out.png` (train loss).
- **Expected Results:** After 100 epochs, test accuracy should be $\geq 65\%$, train accuracy should be $\approx 100\%$, and training loss ≤ 0.3 .



(a) Train/test accuracy



(b) Training loss

Solution:

Test Accuracy: **0.667200**

Train Accuracy: **0.999980**

Training Loss: **0.036354**

Collaboration Survey Please answer the following:

1. Did you receive any help whatsoever from anyone in solving this assignment?

☒ Yes

☐ No

- If you answered 'Yes', give full details:
- (e.g. "Jane Doe explained to me what is asked in Question 3.4")

Discussed with Srecharan about the general approach and methodology for each question

2. Did you give any help whatsoever to anyone in solving this assignment?

☒ Yes

☐ No

- If you answered 'Yes', give full details:
- (e.g. "I pointed Joe Smith to section 2.3 since he didn't know how to proceed with Question 2")

Discussed with Srecharan about the general approach and methodology for each question

3. Note that copying code or writeup even from a collaborator or anywhere on the internet violates the [Academic Integrity Code of Conduct](#).